CNNs vs U-Net: Understanding Deep Learning for Vision

From Image Classification to Pixel-Perfect Segmentation

Made with GAMMA

What is a CNN?

Convolutional Neural Networks learn to recognize **WHAT** is in an image. Small filters (kernels) slide across the image, detecting local patterns at increasing levels of abstraction:

1. Edges
2. Textures
3. Shapes
4. Objects

Pooling Layers

Shrink spatial dimensions, retaining strongest activations for efficiency.

Deep Representations

Deeper layers build increasingly abstract, semantic feature maps.

Classification

Perfect for answering: *"Is this a cat or a dog?"*

Made with GAMMA

The CNN Pipeline

As the image passes through the network, spatial resolution is progressively compressed while semantic understanding grows.

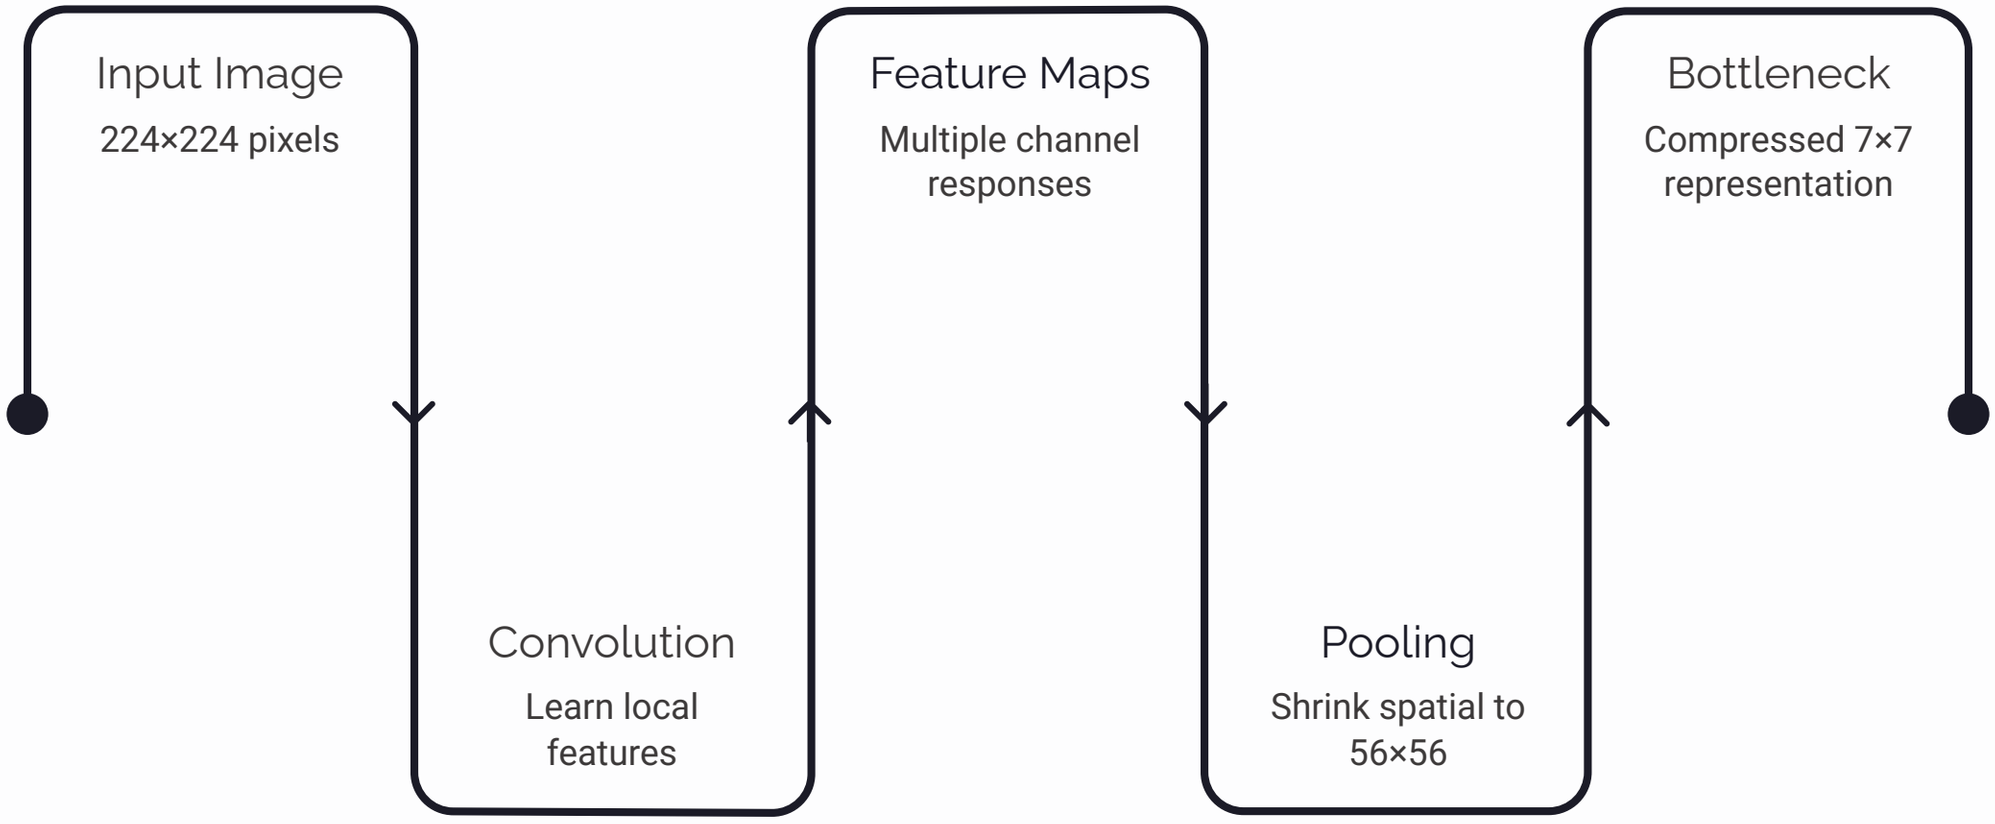

Key Limitation: The network learns **WHAT** is in the image — but forgets **WHERE**. Spatial information is irreversibly compressed by the bottleneck stage.

Made with GAMMA

The Problem CNNs Couldn't Solve

SEMANTIC SEGMENTATION

In medical imaging, knowing **"there is a tumor"** is not enough. Clinicians need to know **exactly which pixels** belong to the tumor — its precise boundaries, shape, and extent.

Semantic segmentation means assigning a class label to every single pixel in the image.

Pooling is the Villain

Aggressively shrinks spatial resolution — 224×224 collapses to 7×7.

Location Info is Lost

Precise pixel positions are permanently destroyed at the bottleneck.

Naive Fix Fails

Removing pooling sacrifices the receptive field needed for global context.

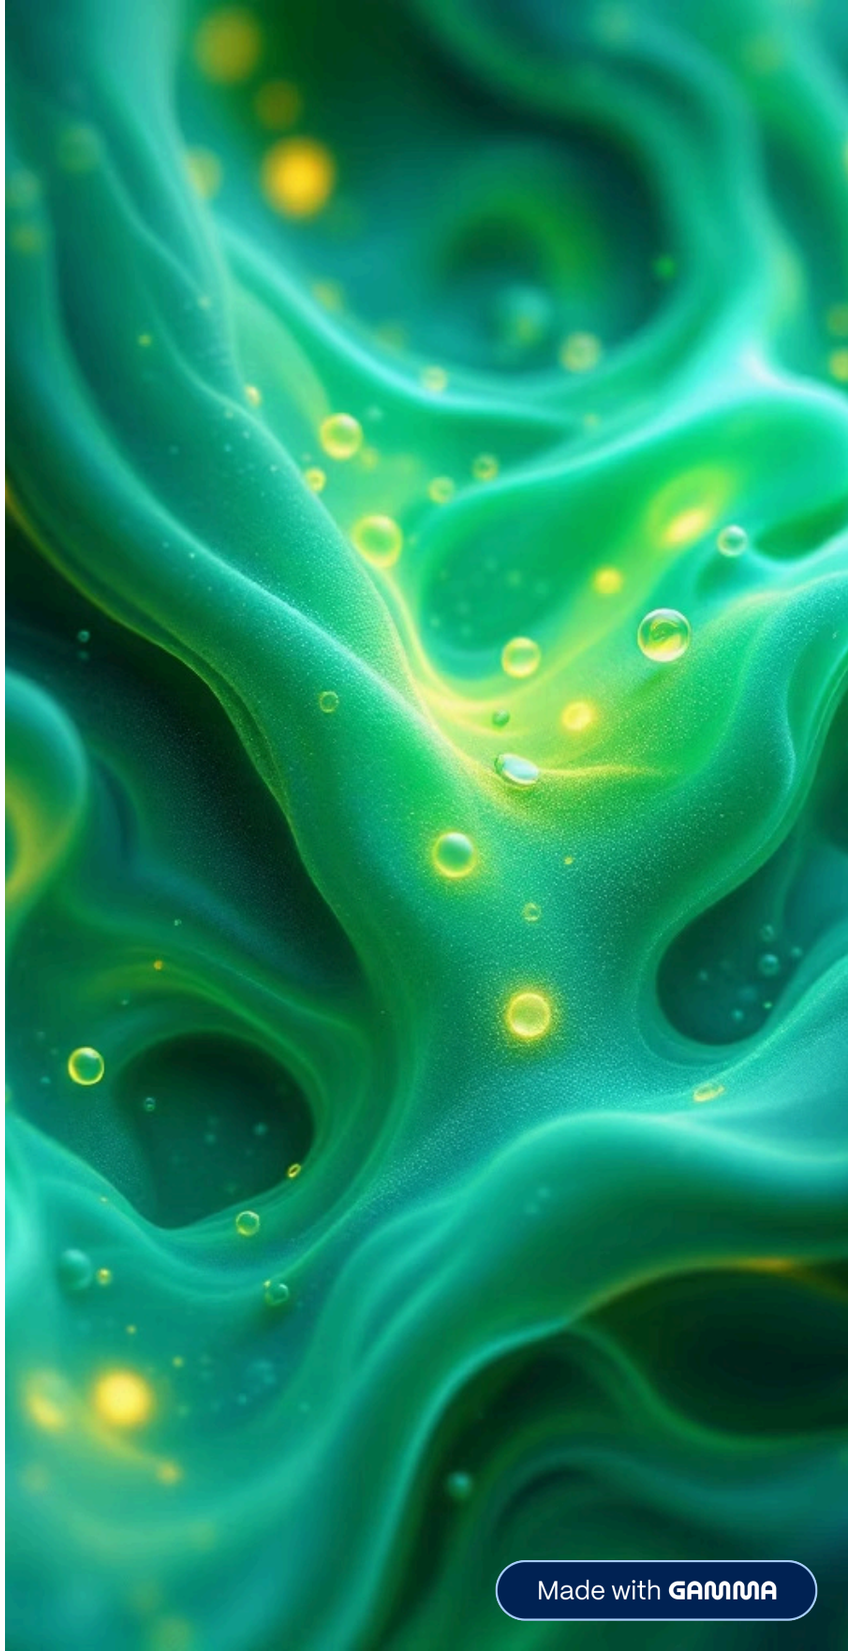

Made with GAMMA

The Core Tension

CNNs face a fundamental architectural contradiction that cannot be resolved within the standard pipeline:

⬇️ Downsample

Build a wide receptive field to understand global **context** — what objects are present and how they relate.

📍 Preserve Precision

Maintain **spatial resolution** to know the exact pixel-level location of each object or boundary.

> A standard CNN can understand **what** something is, or **where** it is — but not simultaneously at the same resolution. These two goals are in direct conflict.

Made with GAMMA

U-Net — The Elegant Solution

RONNEBERGER ET AL., 2015

Introduced for **biomedical image segmentation**, U-Net resolves the core tension with one powerful insight:

> 1. ***Downsample*** *to understand context*
> 2. *Then* ***upsample*** *to recover spatial precision using* ***skip connections*** *to bring back the details lost during downsampling.*

Encoder (Left Arm)

Compresses the image, builds rich semantic context through conv + pooling.

Decoder (Right Arm)

Progressively expands, recovering full spatial resolution using extracted information and feature maps of the encoder.

Skip Connections

Horizontal bridges that transfer precise spatial details directly to the decoder.

Made with GAMMA

The U-Net Architecture

The architecture forms a symmetric **U shape** — an encoder that compresses, a bottleneck of maximum context, and a decoder that restores resolution step by step.

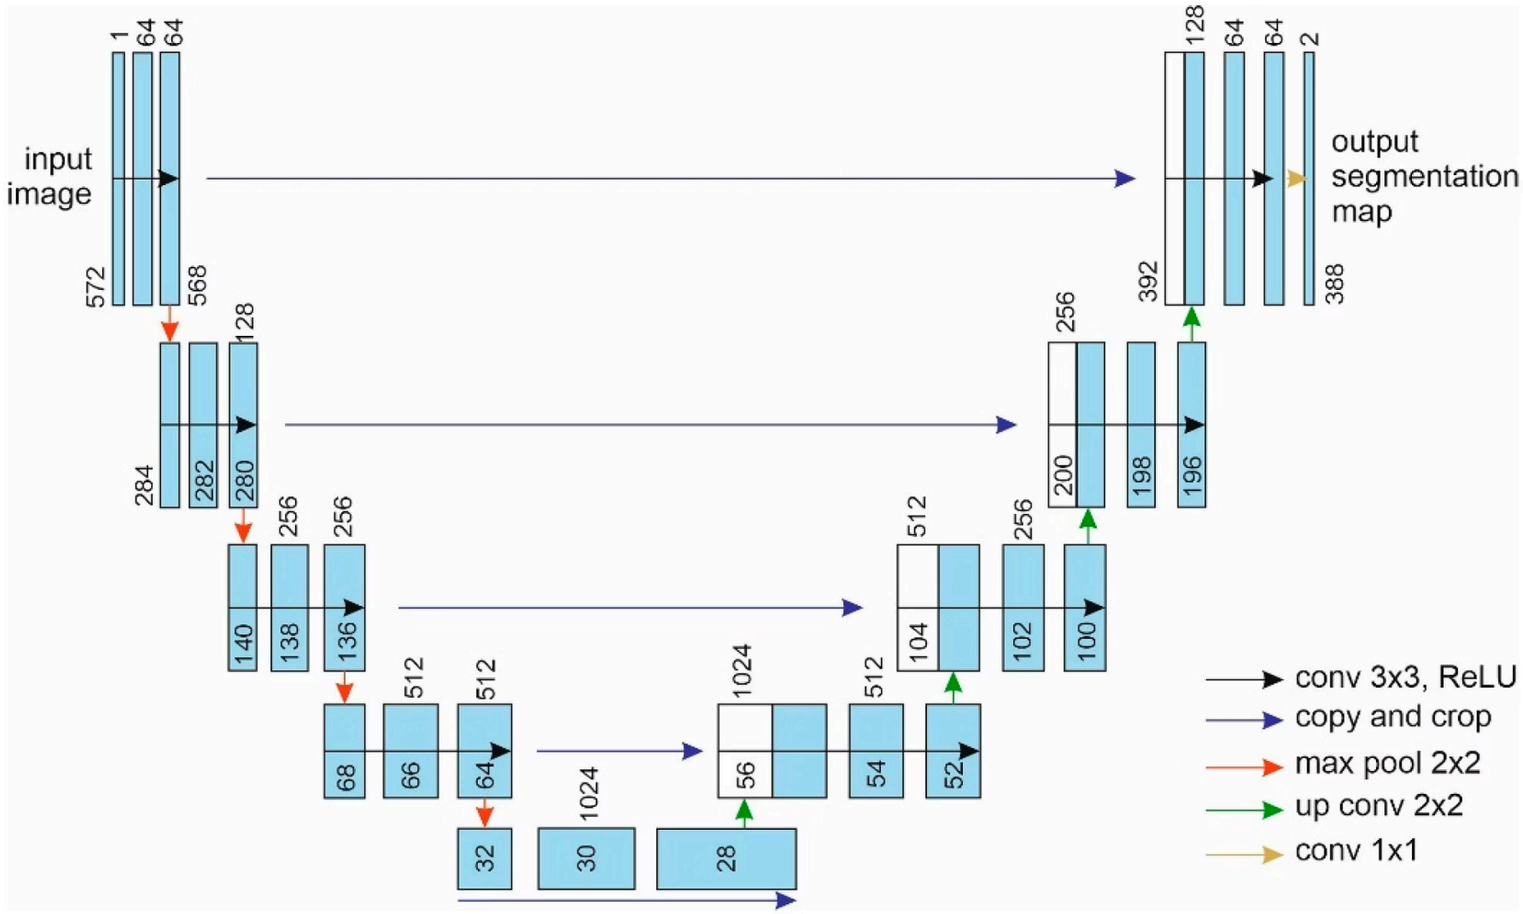

Skip Connections: At every level, encoder feature maps are **concatenated** onto the decoder's upsampled maps — injecting full-resolution spatial detail back into the reconstruction path.

Made with GAMMA

Skip Connections — The Key Innovation

What Happens at Each Level

The encoder's feature map is **concatenated** onto the decoder's upsampled feature map at the matching resolution. The decoder receives two streams simultaneously:

- Deep semantic context from the bottleneck — *"what is this?"*
- Precise spatial detail from the encoder — *"exactly where is it?"*

Without Skip Connections

Upsampling alone produces blurry, imprecise object boundaries — lost detail cannot be reconstructed.

With Skip Connections

The decoder sharpens predictions using full-resolution detail — crisp, accurate boundaries emerge.

The Shortcut Analogy

"Here's exactly what this region looked like at full resolution — use it to reconstruct precisely."

Made with GAMMA

CNN vs U-Net — Side by Side

Dimension	CNN	U-Net
Goal	Image classification	Pixel-level labeling
Output	Single class vector	Full-resolution segmentation mask
Spatial Info	Progressively discarded	Preserved via skip connections
Context	Global via pooling	Global + local simultaneously
Best For	"What is this?"	"Where exactly is each thing?"

CNN Strength

Efficient, powerful classification with compact feature representations

U-Net Strength

Dense prediction preserving spatial fidelity across every pixel

Made with GAMMA

U-Net's Impact and Legacy

Originally designed for biomedical segmentation — where labeled data is scarce — U-Net enabled precise cell boundary detection with very few training images. Its influence now spans all of deep learning.

Dense Prediction

Semantic segmentation, depth estimation, and super-resolution all build on the U-Net paradigm.

Diffusion Models

The denoising step inside Stable Diffusion and DALL-E uses a U-Net at its core.

Universal Pattern

The encoder-decoder with skip connections has become a foundational design pattern across deep learning.

Made with GAMMA